

When to Intervene: Comparing Constant and Physiology-Triggered Robotic Assistance

Anonymous Author(s)

Submission Id: 2613

Abstract

In human–robot interaction (HRI), deciding when to provide assistance requires understanding how different intervention strategies affect both task performance and the user experience. Previous work has compared physiology-triggered robotic assistance with periodic support in surgical training. We extend this work to physical spatial problem solving through a within-participant experiment comparing three conditions: no assistance, constant assistance every 30 seconds, and adaptive assistance triggered by a short-term heart-rate rise. Twenty-four participants completed nine physical tangram attempts across three blocks. Condition orders were balanced across participants, and puzzles were randomly assigned to blocks. Tangrams combine spatial reasoning with physical manipulation, providing a simple task for studying informational and physical assistance. Participants received on-screen hints and robotic-arm interventions through a Wizard-of-Oz setup, with the assigned condition determining intervention timing and task progress guiding assistance content. We evaluated puzzle completion, correct-piece progress, recorded round duration, workload, helpfulness, intervention timing, frustration, and study-wide trust. Both assistance conditions improved completion and correct-piece progress and reduced workload compared with no assistance (Holm-adjusted $p < .001$); both also reduced frustration. Descriptively, constant assistance achieved higher completion and correct-piece progress, while adaptive assistance received higher helpfulness and timing ratings and lower workload and frustration. Together, these findings demonstrate benefits of assistance in physical spatial problem solving and suggest potential trade-offs between task performance and user experience that warrant further investigation.

Keywords

Human–robot interaction, adaptive assistance, intervention timing, physiological sensing, Wizard-of-Oz, tangram

ACM Reference Format:

Anonymous Author(s). 2027. When to Intervene: Comparing Constant and Physiology-Triggered Robotic Assistance. In *Proc. of the 26th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2027)*, Hanoi, Vietnam, 3–7 May 2027, IFAAMAS, 8 pages.

1 Introduction

A pause during a physical puzzle can reflect productive mental rotation or uncertainty about the next move. A helpful intervention

in the latter case may interrupt reasoning in the former. Assistance timing therefore matters alongside assistance content.

Recent systems learn assistance type, timing, and confidence [1] or estimate action completion for proactive coordination [10]. A systematic review of 124 publications identifies timing and role allocation as dimensions of adaptive HRI [13]. Physiological information offers another timing cue: Yang et al. compared workload-adaptive robotic suction with periodic support in surgical training [33]. We extend that comparison to physical spatial problem solving, where text hints guide interpretation and an arm retrieves relevant pieces.

The usefulness of an intervention depends on both its task benefit and its cost to the recipient. Assistance-appropriateness judgments reflect this balance [21], and a proactive escape-room robot elicited more interaction without significantly improving overall performance [31]. Evaluating a timing policy therefore requires task outcomes alongside the experience of receiving help, rather than treating more interaction as an improvement in itself.

Physiological triggering also needs task-specific evaluation. Reviews report heterogeneous workload measures and mixed cardiac findings in human–robot collaboration [18], while cardiac measurement guidelines distinguish heart rate from heart-rate variability and emphasize signal quality and physical activity [20]. We use a short-term heart-rate rise as an intervention cue, not as a direct measure of frustration, and evaluate its consequences through task performance and participant ratings.

Tangrams combine spatial reasoning and manipulation using common materials, observable outcomes, and repeated targets. Related work studies physical tangrams in HRI [24] and in non-robotic cognitive-load measurement [16]. Each participant attempted three puzzles in each condition. We ask **RQ1**: How do the conditions differ in correct-piece progress and completion? **RQ2**: How do they differ in reported workload and frustration, and how do the assisted policies compare in intervention experience? Recorded round durations and final-survey ratings provide descriptive context.

Our contributions are:

- (1) A comparison of periodic and physiology-triggered assistance in physical spatial problem solving, using a within-participant design, an unassisted comparator, and a common hint and retrieval format.
- (2) Participant-level evidence that both assistance policies improve progress and completion and reduce reported workload and frustration relative to no assistance, without establishing superiority of either assisted policy.
- (3) An assessment of intervention experience and users’ subjective feedback on hint frequency and robotic piece retrieval.



This work is licensed under a Creative Commons Attribution International 4.0 License.

2 Related Work

2.1 Assistance timing and proactive coordination

Gaze-based assembly research studies readiness for joint action [15], while LLM-based Attentive Support uses scene and dialogue information to select help or silence, evaluated in constructed scenarios and a robot demonstration [26]. A mentalising/Q-learning memory-game study reports performance and acceptance benefits, but explanation detail differs and timing remains fixed [2]. These studies motivate contextual assistance without isolating physiological versus periodic triggering.

Timing also affects feedback and interruption handling. Reminders before robot behavior changes elicit faster, more immediate feedback in Space Invaders, but not a significant whole-game timing effect [6]. Exploratory analyses of user-initiated interruptions associate unsuccessful handling with lower perceived inclusion and discussion satisfaction [7]. Initiating assistance, soliciting feedback, and responding to interruptions are distinct decisions; our comparison concerns assistance initiation.

2.2 Physiological information in assistance

Yang et al.’s surgical system uses EEG and eye tracking to inform adaptive suction, with a periodic comparator selected to approximate earlier observed assistance frequency [33]. Teo et al. use individualized physiological workload markers to trigger aid during robot supervision, imposing aid later if no trigger occurs [27]. Both inform our timing comparison, although the tasks, sensors, and assistance differ.

Physiological measurement does not necessarily drive online adaptation. Hostettler et al. adapt to user distance while measuring pupil responses [12]; Ojsteršek et al. personalize robot parameters using a skills test and analyze ECG afterward [17]. Korivand et al. report that wristband data could not be integrated directly in real time [14]. Prajod et al. evaluate an HRV-based challenge classifier and separately analyze facial expressions; their adaptive assembly condition uses Wizard-of-Oz progress judgments rather than physiological triggers [19].

Beyond robotic execution, GuideAI combines cardiac, gaze, and behavioral information to adapt learning content, pacing, and feedback [22]. A single-surgeon simulation study identifies subjective workload and mean heart rate among influential task-performance prediction features [32]. Generalizing such signals remains challenging: Bhagat Smith and Adams review portability, model complexity, and adaptability across unknown tasks [4], and Capponi et al. find no clear RMSSD pattern across assembly configurations [8]. Our study evaluates an HR-rise timing policy rather than a workload-prediction model.

2.3 Physical tasks and participant experience

Tangram research examines gaze around robot failures [24], physical-piece assembly with screen-presented targets and wearable sensing in SensCogAR [16], and autonomous robotic assembly from silhouettes in MRChaos [34]. SensCogAR measures cognitive load without robotic assistance, whereas MRChaos addresses robot reasoning and manipulation rather than human assistance

outcomes. A preliminary three-person assembly study compares manual, collaborative, and guided conditions using EEG [5].

Cavicchi et al. identify potential workload and performance costs in shared tasks and benefits of appropriate task-relevant social cues [9]. Human-led assembly collaboration lowers five workload dimensions, while slower robot pacing lowers mental and temporal demand through action-onset timing rather than movement speed [29]. Older adults report greater workload with robot than human guidance in a modified Trail Making Test; the younger-adult contrast is nonsignificant [30]. Workload definitions also differ: simulated order-picking research represents physical workload fairness through lifted mass [23].

In a medical-training simulation, verbal object-search guidance with visual reminders from a Wizard-of-Oz robotic crash cart yields lower workload and higher usefulness and ease-of-use ratings than reversed cues or no feedback [25]. We retain a common assistance format across policies, without assuming equal numbers of hints or retrievals.

2.4 Wizard-of-Oz evaluation

Wizard-of-Oz methods enable human-operated delivery of robot interactions. Thunberg et al.’s workshop proposal emphasizes the practical, ethical, and methodological tensions of the wizard’s role [28]. Interviews with six HRI researchers identify operator response processing, robot delays, unpredictable participants, and control precision as challenges [3]. Our apparatus separates the timing trigger from the researcher’s selection and delivery of task-relevant assistance.

3 Formative Assessment of Tangram Solving

3.1 Assessment procedure and task experience

A formative assessment collected 50 substantive responses: 38 undergraduates, nine graduate students, and three faculty members. Ages ranged from 18 to 30 years ($M = 21.96$, $SD = 3.57$); 25 respondents identified as men, 20 as women, and five preferred not to disclose gender. A Google Form shared through college mailing groups linked to an online square tangram with a timer. Respondents attempted it and reported their experience and expected assistance preferences. Some also participated in the robot experiment; the samples are not treated as independent replications.

Twenty-nine had previously solved tangrams, including four regular solvers; ten had heard of them without solving them and eleven were unfamiliar. Twenty-eight reported complete solutions (56%), 18 partial solutions (36%), and four no solution (8%). Thirty-seven (74%) felt stuck at least once. These self-reports indicate that the task can elicit difficulty even among respondents familiar with tangrams.

3.2 Assistance expectations and task rationale

Respondents independently rated hypothetical assistance from 1 (most helpful) to 5 (least helpful), rather than ranking the options. Mean ratings were 2.58 ($SD = 1.34$) for a simpler puzzle, 2.76 ($SD = 1.22$) for a text hint, 2.80 ($SD = 1.14$) for a video tutorial, and 2.94 ($SD = 1.27$) for a robot showing piece placement (Figure 1). Responses requested first-piece hints, visual placement guidance,

and help at difficult moments. These descriptive expectations motivate the task and assistance modalities; they do not demonstrate robot efficacy.




Figure 1: Expected helpfulness in the formative assessment ($N = 50$). Dots show means and bars show SD; lower 1–5 ratings indicate greater expected helpfulness.

Each experimental tangram required seven colored pieces: two large triangles, one medium triangle, two small triangles, a square, and a parallelogram, together forming a 10×10-inch square. Participants saw nine black silhouette targets on a light-gray background, all in the same format. Colored solutions were available only to the researchers.

Participants interpret the target, rotate and place pieces, and revise arrangements without a prescribed sequence. Hints address orientation or adjacency; retrieval brings a relevant piece to the workspace. Common materials support repeated comparisons across targets.

4 Study Method: Assistance System and Experimental Protocol

4.1 Architecture and apparatus

Our system coordinated three browser interfaces through a local server (Figure 2). The researcher dashboard displayed the colored solution, workspace camera view, timer, cardiac feedback, and assistance controls. The participant laptop presented instructions, silhouette targets, text hints, the timer, and questionnaires. The robot-operator display identified the piece-specific program to execute. The server synchronized the interfaces and recorded task outcomes and questionnaire responses.

The apparatus combined an Orangewood OWL63 arm, a Maxim H Band wrist sensor, and a Logitech C270 workspace camera (Figure 3). An opaque partition separated the participant from two administrators. A camera above the partition covered the table, which contained the assembly area, a drop spot, a laptop, and marked piece-pickup positions. One of seven piece-specific arm programs retrieved the relevant piece to the drop spot; the participant then placed it in the arrangement. Kill switches were accessible to the operator and beside the participant.

At each assistance cue, an administrator selected one preset hint about location, orientation, or adjacency based on the current arrangement. If the relevant piece remained at its marked pickup position, the arm retrieved it and the hint was displayed; otherwise only the hint was displayed. The pickup positions were explained in the instructions. Visual and audible reminders prompted delivery. Assistance was delivered rather than offered for acceptance: participants could ignore the guidance but could not decline its delivery.

4.2 Physiological monitoring

Participants rested before a separate 60-second sensor-calibration recording. Cardiac measurements streamed over Bluetooth to a Python collector. The detector compared median heart rate over the latest 15 seconds with the preceding 60 seconds and raised an HR-rise flag at a difference of at least 5 beats per minute. The calibrated baseline provided the initial reference until the rolling reference was established. Signal-quality checks excluded out-of-range readings and poor sensor contact; sufficient usable samples were required to raise a flag.

The dashboard displays the flag alongside the workspace view. In the adaptive condition, the flag initiates assistance, while the puzzle arrangement determines the content of the hint or robot action. The flag is a heart-rate-rise cue for assistance timing; participant workload and frustration are assessed through questionnaires.

4.3 Assistance conditions

No assistance (control) provided no hints or robot interventions. **Constant assistance** delivered help every 30 seconds. **Adaptive assistance** delivered help when the HR-rise flag was raised, with a 15-second wait before help for rapid or sustained recurrence. Both assisted conditions used the same preset-hint list and pickup-availability rule. The condition determined timing; task state determined the next hint and whether retrieval accompanied it. This compares assistance policies, not equal-dose or robot-only interventions.

4.4 Participants and procedure

Twenty-four volunteers recruited through forms shared in college groups received coupons. Most were undergraduates; graduate students, teaching fellows, and research fellows also participated. Ages ranged from 20 to 27 years ($M = 21.96$, $SD = 2.01$); 14 were men and ten women. Consent and background questions preceded the task.

Each participant attempted nine distinct puzzles in three three-puzzle blocks, with one condition per block. Four participants followed each of the six condition orders; puzzles were randomly assigned to blocks. Balanced condition orders did not produce identical puzzle allocation across conditions. The supplementary material provides the order diagram and allocation counts.

Rounds ended when the target was solved or at the five-minute beep, when participants stopped working. The protocol included a questionnaire after each round, five-minute breaks between blocks, and a final survey after round nine. Participants were told about the hidden operator after completing the experiment.




Figure 2: Assistance architecture. Cardiac measurements pass through signal-quality checks and a heart-rate-rise detector. The assigned condition determines timing; task context guides hint selection and pickup availability determines retrieval. The local server synchronizes interfaces and records study measures.

4.5 Measures and analysis

Task outcomes were binary completion and correctly placed pieces, recorded in researcher notes and photographs of final arrangements. Placement and orientation were judged by humans against the target solution. Counts ranged from 0 to 7; a solved arrangement used all seven pieces correctly. Where multiple orientations were admissible, scoring used the orientation yielding the highest correct-piece count. Recorded round duration included solved and unsolved attempts and is reported descriptively, not as time to successful completion or verified active solving time.

Round-specific questionnaires assess six adapted NASA-TLX items [11] on seven-point scales (1–7): mental, physical, and temporal demand, performance, effort, and frustration. Performance runs from failure to perfect and is reversed as $8 - P$. We average the six aligned burden ratings without weighting and rescale the mean B as $(B - 1)100/6$. Assisted rounds also assess helpfulness, timing, seamlessness, and frustration relief. Final helpfulness, efficacy, trust, and reliance ratings summarize the whole study, not individual conditions.

All 24 participants contributed task outcomes, round ratings, and final surveys, with 72 attempts per condition. Two-sided paired t tests compared each participant’s three-round condition averages. We report paired mean differences, pointwise 95% confidence intervals, and Cohen’s d_z .

Holm correction limits false-positive findings across related comparisons. The main group contains twelve tests: three condition contrasts for each of completion, correct pieces, workload, and frustration. A separate nineteen-test group covers the other five workload items and four adaptive–constant assistance ratings. Throughout, p_H denotes the Holm-adjusted value; significance uses $p_H < .05$. These analyses are exploratory. A supplementary puzzle- and block-period-adjusted robustness check supported the same main contrast conclusions.

5 Findings

Descriptive means favored constant assistance on objective completion and correct-piece progress, and adaptive assistance on subjective overall workload, frustration, and intervention experience. Corrected comparisons did not establish a difference between the assisted policies.

5.1 Task performance

Both assisted conditions improved task performance relative to control (Table 1; Figure 4). Across 216 attempts, 90 puzzles were completed: 15.3% in control, 59.7% with constant assistance, and 50.0% with adaptive assistance.




Figure 3: Wizard-of-Oz setup: schematic (left); robot and workspace, operator interface, and overhead view (right, top to bottom). Images are cropped and identifying details redacted.

Table 1: Performance ($n = 24$): participant-condition means (SDs) for pieces and duration; completion totals across 72 attempts per condition.

Measure	Control	Constant	Adaptive
Pieces / 7	3.43 (1.40)	5.68 (1.05)	5.43 (1.13)
Completed / 72	11	43	36
Duration (s)	292.4 (22.5)	248.2 (53.3)	265.4 (29.1)

Relative to control, constant and adaptive assistance increased correct-piece counts by 2.25 and 2.00 pieces and completion by 44.4 and 34.7 percentage points, respectively (all $p_H < .001$). Constant had higher observed progress and completion than adaptive, but those contrasts were not significant ($p_H = .532$ and $.444$, respectively).

Mean recorded round durations were 292.4, 248.2, and 265.4 seconds for control, constant, and adaptive assistance. These descriptive values combine successful and unsuccessful attempts and include end-marking time; they do not establish solving-speed differences.

5.2 Subjective workload

Adapted workload averaged 62.46 in control, 49.50 with constant assistance, and 46.06 with adaptive assistance (Table 2; Figure 4). Reductions against control were 12.96 points for constant (95% CI [8.52, 17.40], $t(23) = -6.04$, $d_z = -1.23$) and 16.40 for adaptive (CI [10.84, 21.95], $t(23) = -6.10$, $d_z = -1.25$), both $p_H < .001$. The adaptive-constant contrast was not significant ($p_H = .532$).

Table 2: Workload means (SDs), $n = 24$. Items use 1–7: higher performance is better; higher demand, effort, and frustration indicate greater burden. Overall workload uses 0–100.

Measure	Control	Constant	Adaptive
Mental demand	5.71 (0.94)	4.82 (0.98)	4.63 (1.46)
Physical demand	3.65 (1.68)	3.28 (1.37)	3.06 (1.60)
Temporal demand	3.92 (1.38)	3.63 (1.25)	3.56 (1.45)
Performance	2.93 (1.39)	4.54 (1.47)	4.94 (1.23)
Effort	5.44 (0.69)	4.71 (0.95)	4.78 (1.19)
Frustration	4.69 (1.32)	3.93 (1.23)	3.51 (1.06)
Workload /100	62.46 (11.26)	49.50 (12.93)	46.06 (15.28)

Frustration averaged 4.69 in control, 3.93 with constant assistance, and 3.51 with adaptive assistance. Both assisted conditions reduced frustration versus control: constant by 0.76 points (95% CI [0.25, 1.28], $p_H = .027$) and adaptive by 1.18 points (CI [0.66, 1.70], $p_H < .001$). Adaptive minus constant was -0.42 (CI $[-0.90, 0.07]$, $t(23) = -1.77$, $p_H = .363$; Figure 5), which did not establish a policy difference.

Among secondary ratings, both assisted conditions reduced mental demand and increased perceived performance versus control (all $p_H < .001$). Constant reduced effort ($p_H = .003$); adaptive's effort contrast was nonsignificant ($p_H = .064$). Physical and temporal demand comparisons were nonsignificant. No item differed significantly between assisted policies.




Figure 4: Paired contrasts ($n = 24$): mean differences and pointwise 95% CIs. Labels show $t(23)$ and p_H from the twelve-test main group. K: constant; C: control; A: adaptive. Blue indicates $p_H < .05$; gray indicates nonsignificance.

5.3 Intervention experience

Adaptive received descriptively higher means on all four assistance items (Table 3; Figure 5). The timing contrast was +0.56 points ($t(23) = 2.91$, unadjusted $p = .008$), but nonsignificant after correction ($p_H = .103$). Helpfulness, seamlessness, and frustration-relief comparisons also did not pass correction. These ratings concern the combined experience of hints and retrieval, not either modality separately.

Table 3: Assistance ratings, mean (SD), $n = 24$. All items use 1–7, with higher ratings indicating a better experience.

Measure	Constant	Adaptive
Helpfulness	4.58 (1.26)	4.83 (1.35)
Timing	4.68 (1.48)	5.24 (1.40)
Seamlessness	4.01 (1.39)	4.10 (1.39)
Frustration relief	4.42 (1.35)	4.82 (1.46)

5.4 Overall experience

Across 24 final surveys, mean helpfulness was 5.00 ($SD = 1.53$), efficacy 4.83 ($SD = 1.55$), trust 4.79 ($SD = 1.59$), and agreement with following uncertain robot guidance 4.88 ($SD = 1.48$), all on 1–7 scales. These ratings summarize the overall study experience rather than individual conditions.

Two final-survey comments illustrate the experience of assistance. One participant wrote, “The hints were helpful and overall well timed, but the actual robot arm moving the pieces [was] not adding anything more.” Another wrote, “Sometimes hints being given too often was counterintuitive and more distracting [than] better-spaced hints, even though any hint helps.” These comments distinguish guidance from the physical contribution of retrieval and highlight the costs of frequent intervention.




Figure 5: Constant-adaptive comparison ($n = 24$), means $\pm$ SD. Lower frustration is better; higher assistance ratings are better. Labels show paired-test p_H ; none passes .05. SD bars describe participant variability, not uncertainty in the paired difference.

6 Discussion, Limitations, and Future Work

6.1 Task progress and the experience of assistance

Both assisted conditions improved progress and completion and reduced reported workload and frustration relative to control. Descriptively, constant had higher completion and correct-piece counts, while adaptive had lower overall workload and frustration and higher assistance ratings. However, no corrected assisted-policy comparison established a difference. A practical choice between the policies remains unresolved; nonsignificance does not show equivalence.

Regular hints may provide repeated opportunities to advance an arrangement, whereas HR-rise-triggered hints may concentrate help at different moments. These are explanations to test, not mechanisms established by the present data. The final-survey comments suggest evaluating content and physical contribution separately: piece retrieval leaves placement to the participant and may add movement overhead, while frequent hints may interrupt reasoning.

6.2 Physiological feedback and limitations

The trigger detects an HR rise, not validated psychological stress. Wearable cardiac measurements are sensitive to movement and sensor contact [20]. The observed assistance benefits therefore do not establish that the HR-rise flag measures stress or workload. Researchers observed some waiting for assistance, but did not systematically code its prevalence or condition distribution. Participants were not told the adaptive trigger rules. Future studies could test explanations and user-requested help, while accounting for PACE's finding that explicit queries can increase waiting [10].

The sample comprised 24 college-recruited volunteers aged 20–27 years, mostly undergraduates, with graduate students, teaching fellows, and research fellows. This recruitment context limits generalization to older adults and experienced industrial workers. A supplementary analysis accounting for puzzle and block period supported the main conclusions; it does not remove every possible learning or deployment effect. Completion and correct-piece counts are linked, and frustration contributes to the workload composite, so these findings are not independent replications.

The experiment evaluates assistance packages: task state determines whether retrieval accompanies a hint, and the policies need not deliver equal amounts of help. It therefore does not isolate timing from assistance dose or establish the added benefit of the arm. Future work should compare hints with retrieval or direct placement separately and relate help to observable progress. Tangrams share orientation and placement demands with assembly [34], while retrieval resembles material provision in order picking [23]; transfer requires separate safety, precision, and workflow evaluation.

7 Ethical Considerations

Participation was voluntary, with prior consent and the option to pause or stop without penalty. Data included cardiac measurements, questionnaires, and consented workspace viewing

or recording. Participants were debriefed about the concealed Wizard-of-Oz operator after the experiment.

8 Conclusion

We compared unassisted, periodic, and HR-rise-triggered assistance in 24 participants' physical tangram attempts. Both assisted conditions improved correct-piece progress and completion and reduced adapted workload and frustration relative to control. Descriptively, objective completion and correct-piece progress favored constant assistance, while subjective workload, frustration, and assistance ratings favored adaptive assistance. No corrected contrast established a difference between the assisted policies, including in the puzzle/period-adjusted robustness analysis. The results support assistance in this task while leaving policy choice, physiological validity, and the added value of robot retrieval open.

AI Assistance Disclosure

AI-assisted tools supported manuscript preparation and analysis-code development. They were not used to formulate hypotheses or design or conduct the experiment. The authors retain responsibility for the analyses, text, and citations.

References

- [1] Antonio Andriella, Ilenia Cucciniello, Antonio Origlia, and Silvia Rossi. 2025. A Bayesian framework for learning proactive robot behaviour in assistive tasks. *User Modeling and User-Adapted Interaction* 35, 1, Article 1 (2025). doi:10.1007/s11257-024-09421-1
- [2] Antonio Andriella, Giovanni Falcone, and Silvia Rossi. 2025. Enhancing Robot Assistive Behaviour by Mentalising User Intent and Beliefs with Reinforcement Learning. *International Journal of Social Robotics* 17 (2025), 2293–2309. doi:10.1007/s12369-025-01280-z
- [3] Alexandra Bejarano, Saad Elbeleidy, Terran Mott, Sebastian Negrete-Alamillo, Luis Angel Armenta, and Tom Williams. 2024. Hardships in the Land of Oz: Robot Control Challenges Faced by HRI Researchers and Real-World Teleoperators. In *2024 IEEE 33rd International Conference on Robot and Human Interactive Communication (RO-MAN)*. 1914–1921. doi:10.1109/RO-MAN60168.2024.10731251
- [4] Joshua Bhagat Smith and Julie A. Adams. 2026. A survey of machine learning for estimating workload: considering unknown tasks. *Frontiers in Robotics and AI* 13, Article 1872363 (2026). doi:10.3389/frobt.2026.1872363
- [5] Carlo Caiazzo, Marija Savkovic, Milos Pusica, Nastasija Nikolic, Ivan Macuzic, and Marko Djapan. 2024. Comparative Analysis Of Mental Workload In Adaptive Human-Robot Collaboration During Assembly Tasks. In *Advances in Reliability, Safety and Security, Part 5 (ESREL 2024)*. <https://scidar.kg.ac.rs/handle/123456789/21811>
- [6] Kate Candon, Helen Zhou, Sarah Gillet, and Marynel Vázquez. 2023. Verbally Soliciting Human Feedback in Continuous Human-Robot Collaboration: Effects of the Framing and Timing of Reminders. In *Proceedings of the 2023 ACM/IEEE International Conference on Human-Robot Interaction*. 290–300. doi:10.1145/3568162.3576980
- [7] Shiye Cao, Jiwon Moon, Amama Mahmood, Victor Nikhil Antony, Ziang Xiao, Anqi Liu, and Chien-Ming Huang. 2025. Interruption Handling for Conversational Robots. In *Proceedings of Robotics: Science and Systems XXI*. doi:10.15607/RSS.2025.XXI.089
- [8] Matteo Capponi, Riccardo Gervasi, Luca Mastrogiacomo, and Fiorenzo Franceschini. 2024. Assembly complexity and physiological response in human-robot collaboration: Insights from a preliminary experimental analysis. *Robotics and Computer-Integrated Manufacturing* 89, Article 102789 (2024). doi:10.1016/j.rcim.2024.102789
- [9] Shari Cavicchi, Abdulaziz Abubshait, Giulia Siri, Magda Mustile, and Francesca Ciardo. 2025. Can humanoid robots be used as a cognitive offloading tool? *Cognitive Research: Principles and Implications* 10, Article 17 (2025). doi:10.1186/s41235-025-00616-7
- [10] Davide De Lazzari, Matteo Terreran, Giulio Giacomuzzo, Siddarth Jain, Pietro Falco, Ruggero Carli, and Diego Romeres. 2025. PACE: Proactive Assistance in Human-Robot Collaboration through Action-Completion Estimation. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, 6725–6731. doi:10.1109/ICRA55743.2025.11127399

- [11] Sandra G. Hart. 2006. NASA-Task Load Index (NASA-TLX); 20 Years Later. *Proceedings of the Human Factors and Ergonomics Society Annual Meeting* 50, 9 (2006), 904–908. doi:10.1177/154193120605000909
- [12] Damian Hostettler, Simon Mayer, Jan Liam Albert, Kay Erik Jenss, and Christian Hildebrand. 2025. Real-Time Adaptive Industrial Robots: Improving Safety And Comfort In Human-Robot Collaboration. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. 1–16. doi:10.1145/3706598.3713889
- [13] Bsher Karbouj, Rajwinder Garha, Konstantin Keßler, and Jörg Krüger. 2026. Adaptive Robotic Behavior in Industrial Human–Robot Collaboration: A Systematic Review of Taxonomies, Enabling Mechanisms, and Research Frontiers. *IEEE Access* 14 (2026), 1398–1422. doi:10.1109/access.2025.3649702
- [14] Soroush Korivand, Gustavo Galvani, Arash Ajoudani, Jiaqi Gong, and Nader Jalili. 2024. Optimizing Human–Robot Teaming Performance through Q-Learning-Based Task Load Adjustment and Physiological Data Analysis. *Sensors* 24, 9 (2024), 2817. doi:10.3390/s24092817
- [15] Matteo Lavit Nicora, Pooja Prajod, Marta Mondellini, Giovanni Tauro, Rocco Vertechy, Elisabeth André, and Matteo Malosio. 2024. Gaze detection as a social cue to initiate natural human-robot collaboration in an assembly task. *Frontiers in Robotics and AI* 11, Article 1394379 (2024). doi:10.3389/frobt.2024.1394379
- [16] Javier Melo, Leyla Akinci, Ko Watanabe, Nicolas Großmann, Shoya Ishimaru, and Andreas Dengel. 2026. SensCogAR: Cognitive Load Estimation Via Movement Data in Assembly Tasks. *International Journal of Activity and Behavior Computing* 2026, 1 (2026), 1–20. doi:10.60401/ijabc.147
- [17] Robert Ojsteršek, Borut Buchmeister, and Aljaž Javernik. 2024. Personalizing Human–Robot Workplace Parameters in Human-Centered Manufacturing. *Machines* 12, 8 (2024), 546. doi:10.3390/machines12080546
- [18] Eduarda Pereira, Luis Sigcha, Emanuel Silva, Adriana Sampaio, Nuno Costa, and Nelson Costa. 2025. Capturing Mental Workload Through Physiological Sensors in Human–Robot Collaboration: A Systematic Literature Review. *Applied Sciences* 15, 6 (2025), 3317. doi:10.3390/app15063317
- [19] Pooja Prajod, Matteo Lavit Nicora, Marta Mondellini, Matteo Meregalli Falerni, Rocco Vertechy, Matteo Malosio, and Elisabeth André. 2024. Flow in human-robot collaboration—multimodal analysis and perceived challenge detection in industrial scenarios. *Frontiers in Robotics and AI* 11, Article 1393795 (2024). doi:10.3389/frobt.2024.1393795
- [20] Karen S. Quigley, Peter J. Gianaros, Greg J. Norman, J. Richard Jennings, Gary G. Berntson, and Eco J. C. de Geus. 2024. Publication guidelines for human heart rate and heart rate variability studies in psychophysiology: Part 1: Physiological underpinnings and foundations of measurement. *Psychophysiology* 61, Article e14604 (2024). doi:10.1111/psyp.14604
- [21] Rebecca Ramnauth, Dražen Brščić, and Brian Scassellati. 2026. To Help or Not to Help?: An Expanded Framework for Deciding Socially Appropriate Robot Assistance. *ACM Transactions on Human–Robot Interaction* 15, 3, Article 64 (2026), 32 pages. doi:10.1145/3797264
- [22] Ananya Shukla, Chaitanya Modi, Satvik Bajpai, and Siddharth Siddharth. 2026. GuideAI: A Real-time Personalized Learning Solution with Adaptive Interventions. In *Proceedings of the 31st International Conference on Intelligent User Interfaces*. ACM, 17 pages. doi:10.1145/3742413.3789125
- [23] Igor G. Smit, Zaharah Bukhsh, Mykola Pechenizkiy, Kostas Alogariastos, Kasper Hendriks, and Yingqian Zhang. 2024. Learning Efficient and Fair Policies for Uncertainty-Aware Collaborative Human-Robot Order Picking. arXiv preprint. arXiv:2404.08006 doi:10.48550/arXiv.2404.08006
- [24] Ramtin Tabatabaei, Vassilis Kostakos, and Wafa Johal. 2025. Gazing at Failure: Investigating Human Gaze in Response to Robot Failure in Collaborative Tasks. In *2025 20th ACM/IEEE International Conference on Human–Robot Interaction (HRI)*. 939–948. doi:10.1109/hri61500.2025.10973935
- [25] Tauhid Tanjim, Jonathan St. George, Kevin Ching, and Angelique Taylor. 2025. Help or Hindrance: Understanding the Impact of Robot Communication in Action Teams. In *2025 34th IEEE International Conference on Robot and Human Interactive Communication (RO-MAN)*. 1460–1465. doi:10.1109/RO-MAN63969.2025.11217909
- [26] Daniel Tanneberg, Felix Ocker, Stephan Hasler, Joerg Deigmoeller, Anna Belardinelli, Chao Wang, Heiko Wersing, Bernhard Sendhoff, and Michael Gienger. 2024. To Help or Not to Help: LLM-based Attentive Support for Human-Robot Group Interactions. In *2024 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. 9130–9137. doi:10.1109/IROS58592.2024.10801517
- [27] Grace Teo, Lauren Reinerman-Jones, Gerald Matthews, James Szalma, Florian Jentsch, and Peter Hancock. 2018. Enhancing the effectiveness of human-robot teaming with a closed-loop system. *Applied Ergonomics* 67 (2018), 91–103. First published online 3 October 2017. doi:10.1016/j.apergo.2017.07.007
- [28] Sofia Thunberg, Mafalda Gamboa, Meagan B. Loerakker, Patricia Alves-Oliveira, and Hannah R. M. Pelikan. 2026. Unpacking Lived Experiences of Wizards of Oz. In *Companion Proceedings of the 21st ACM/IEEE International Conference on Human–Robot Interaction*. ACM, 1399–1401. Workshop proposal. doi:10.1145/3776734.3788834
- [29] Wietse van Dijk, Saskia J. Baltrusch, Ezra Dessers, and Michiel P. de Looze. 2023. The effect of human autonomy and robot work pace on perceived workload in human-robot collaborative assembly work. *Frontiers in Robotics and AI* 10, Article 1244656 (2023). doi:10.3389/frobt.2023.1244656
- [30] Simone Varrasi, Roberto Vagnetti, Nicola Camp, John Hough, Alessandro Di Nuovo, Sabrina Castellano, and Daniele Magistro. 2026. Human and Robot Assistance for Cognitive Load in Younger and Older Adults: Multimodal Within-Subject Experimental Study. *Journal of Medical Internet Research* 28, Article e94738 (2026). doi:10.2196/94738
- [31] Thomas Vitry, Vanessa Maeder, Kieran von Valeburg, Asihati Hazaiti, Doga Deniz Ates, Connor Gäde, Jan-Gerrit Habekost, Dennis Becker, and Stefan Wermter. 2026. When May I Help You? On The Effect of Proactivity on Group Human-Robot Collaboration. In *2026 IEEE 35th International Conference on Robot and Human Interactive Communication (RO-MAN)*. arXiv:2606.28469 doi:10.48550/arXiv.2606.28469
- [32] Kaiqi Wei, Chika Kimura, Megumi Shimura, Yoshihiro Shimomura, Xue Zhao, Takaaki Tamura, and Shinichi Sakamoto. 2025. Predicting task performance in robot-assisted surgery using physiological stress and subjective workload: a case study with interpretable machine learning. *Frontiers in Human Neuroscience* 19, Article 1611524 (2025). doi:10.3389/fnhum.2025.1611524
- [33] Jing Yang, Juan Antonio Barragan, Jason Michael Farrow, Chandru P. Sundaram, Juan P. Wachs, and Denny Yu. 2024. An Adaptive Human-Robotic Interaction Architecture for Augmenting Surgery Performance Using Real-Time Workload Sensing—Demonstration of a Semi-autonomous Suction Tool. *Human Factors: The Journal of the Human Factors and Ergonomics Society* 66, 4 (2024), 1081–1102. First published online 11 November 2022; journal issue April 2024. doi:10.1177/00187208221129940
- [34] Chao Zhao, Chunli Jiang, Lifan Luo, Guanlan Zhang, Hongyu Yu, Michael Yu Wang, and Qifeng Chen. 2025. Master Rules from Chaos: Learning to Reason, Plan, and Interact from Chaos for Tangram Assembly. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE. doi:10.1109/ICRA55743.2025.11127610